

# Project Overview

## Contents

|                                 |   |
|---------------------------------|---|
| 1. Workflow .....               | 1 |
| 2. Results .....                | 2 |
| 3. Prompt Engineering .....     | 2 |
| 3.1. Initial input .....        | 2 |
| 3.2. The Factory .....          | 3 |
| The Pitch Picking Process ..... | 4 |
| 4. Automation .....             | 5 |

## 1. Workflow

### ▼ Described by GPT

Proces został zaprojektowany jako szybki, powtarzalny pipeline oparty na podejściu audio-first. Najpierw ręcznie zebrano i uproszczono brief oraz treści ze strony produktu, tworząc jedno źródło prawdy (SSOT), które ogranicza halucynacje modeli i przyspiesza iteracje. Na tej podstawie wygenerowano trzy różne koncepcje kreatywne, z których wybrano jedną najlepiej dopasowaną do grupy docelowej (opiekunowie seniorów) oraz celu kampanii (CTR → strona).

Następnie stworzono skrypt voiceover w języku polskim, zoptymalizowany pod naturalność i synchronizację (pauzy jako punkty cięcia). Do generowania głosu użyto ElevenLabs (model v3), z kontrolą emocji poprzez tagi i iteracyjnym dopracowaniem intonacji.

Na bazie VO opracowano sceny wideo i prompty zoptymalizowane pod Runway, uwzględniając jego ograniczenia (pojedyncza akcja, jawna kamera, minimalizacja dryfu). Generacja odbywała się per-ujęcie, co zwiększa kontrolę i redukuje błędy. Materiały złożono i zsynchronizowano w Adobe After Effects, który pełni rolę warstwy kontroli (timing, korekcja, spójność).

Miniatury wygenerowano osobno, z naciskiem na jasny przekaz i szybkie zrozumienie kontekstu (hook-first). Każdy etap zawierał krótką pętlę iteracyjną: generacja → ocena → poprawa promptu. Całość pipeline'u pozwala przejść od pomysłu do gotowych materiałów w ciągu kilku

## 2. Results

Drive upload of output images and video can be found [here](#).

## 3. Prompt Engineering

Prompts that had an impact on the entire process

### 3.1. Initial input

Warmup and tool verification

This:

vs: Midjourney, Runway, Pika, Leonardo, ElevenLabs, ChatGPT, Claude

Was answered with this:

| Platform    | Core domain | Primary use       | Modality       | Strength               | Weakness              | Control (params) | Custom models | API access | Real-time | Output quality   | Speed  | Best fit             |
|-------------|-------------|-------------------|----------------|------------------------|-----------------------|------------------|---------------|------------|-----------|------------------|--------|----------------------|
| Midjourney  | Image       | Artistic renders  | Img            | Style quality          | No video / API limits | Medium           | No            | Limited    | No        | High (art)       | Medium | Concept art          |
| Runway      | Video       | Video editing/gen | Vid+Img        | End-to-end video       | Cost / complexity     | High             | Yes           | Yes        | Partial   | High (video)     | Medium | Video pipelines      |
| Pika        | Video       | Text→video        | Vid            | Fast gen video         | Less control          | Low-Med          | No            | Limited    | Partial   | Med-High         | Fast   | Quick clips          |
| Leonardo AI | Image       | Asset generation  | Img            | Control + finetune     | Inconsistent          | High             | Yes           | Yes        | No        | High             | Medium | Game assets          |
| ElevenLabs  | Audio       | Voice cloning     | Audio          | Realistic voices       | Narrow scope          | Medium           | Yes (voices)  | Yes        | Yes       | High (voice)     | Fast   | Voice apps           |
| ChatGPT     | Text/MM     | General AI        | Text+Img+Audio | Versatility            | Not top in media gen  | High             | Yes           | Yes        | Yes       | High (text/code) | Fast   | General tasks        |
| Claude      | Text/MM     | General AI        | Text+Img       | Long context reasoning | Smaller ecosystem     | Medium           | Limited       | Yes        | Limited   | High (reasoning) | Medium | Analysis / long docs |

Defining framework that served as a jumping-off point and per-stage branching.

*Redacted.*

With few minor tweaks, this was the proposed pipeline:

| Pipeline overview     |                            |                          |                                                                    |
|-----------------------|----------------------------|--------------------------|--------------------------------------------------------------------|
| Phase                 | Output                     | Prerequisites            | Critical sub-steps                                                 |
| 1. Brief/site intake  | Clean working brief        | PDF + source site access | extract product, target, hooks, brand cues, pitfalls               |
| 2. Concept generation | 3 pitch variants           | phase 1                  | make 3 distinct angles, not 3 rewrites                             |
| 3. Concept selection  | 1 chosen pitch             | phase 2                  | score for fit, clarity, production cost, hook strength             |
| 4. Voiceover script   | Polish VO draft            | chosen pitch             | natural speech, shot-pauses, emotional tone, character distinction |
| 5. Shot script        | per-shot production script | VO draft                 | split by cut, sync points, fallback to per-video if low complexity |
| 6. Generation         | audio + visuals            | shot script              | ElevenLabs first, then visual generation, then sync                |
| 7. Edit/composite     | final masters              | generated assets         | timing, captions, cleanup, loop ending for shorts                  |
| 8. Variant export     | Reel + YouTube + posts     | final master             | adapt aspect ratio, CTA placement, cut length                      |

## 3.2. The Factory

I gave my LLM a condensed brief combined with claims from the test product's website.

From this - with minimal input from my side - GPT has come up with target audience psych eval., tailored 3 pitches for emotionally resonant videos, argued with itself which one's best, iterated on the winner, and finally generated:

### ▼ VO Scripts

- [Local: vo.asc](#)
- this required zero corrections
- Eleven v3 engine
  - the inflection control is very fun to play with
  - control over speaker's wide range of emotions was jaw-dropping

### ▼ VID Scripts

- [Local: vid.asc](#)

- this required quite a few re-iterations on all the video prompts to align with Runway quirks

#### ▼ *IMG Prompts*

- [Local: img.asc](#)
- those were adapted first to Midjourney, re-adapted to Leonardo & GPT
- for the 2<sup>nd</sup> image I had to intervene and redesign the prompt, as it consistently produced very bland results

#### ▼ *CTR Scripts*

- [Local: ctr.asc](#)

## The Pitch Picking Process

For your viewing pleasure

| Quick comparison                                                                                                                                                                                                                                                                                                                    |         |         |           |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------|---------|-----------|
| Metric                                                                                                                                                                                                                                                                                                                              | Pitch 1 | Pitch 2 | Pitch 3   |
| Emotional intensity                                                                                                                                                                                                                                                                                                                 | High    | Medium  | High      |
| Clarity of product                                                                                                                                                                                                                                                                                                                  | Medium  | High    | Medium    |
| Hook strength (<3s)                                                                                                                                                                                                                                                                                                                 | Medium  | High    | High      |
| Production complexity                                                                                                                                                                                                                                                                                                               | Medium  | Low     | Medium    |
| CTR potential                                                                                                                                                                                                                                                                                                                       | High    | Medium  | Very High |
| <b>Recommendation (based on pipeline + constraints)</b><br>→ Pitch 3 ("You think they're safe")<br><br>Reason: <ul style="list-style-type: none"> <li>• aligns with caregiver psychology</li> <li>• strongest hook in &lt;3s</li> <li>• works for both Reel + YouTube</li> <li>• supports all 5 required hooks naturally</li> </ul> |         |         |           |

# 4. Automation

A lot of my 'work' was copy-pasting prompts, switching between LLM's, switching between accounts - annoyances that could be automated away.

There are tools that can help with automation - quite a few of them.

Pipeline Fit Comparison (AI-heavy motion + automation)

| Platform           | Type                  | Role in this project                 | Strength                         | Weakness                            | Fit                  |
|--------------------|-----------------------|--------------------------------------|----------------------------------|-------------------------------------|----------------------|
| Krea (Nodes)       | Visual AI graph       | Rapid ideation (image/video prompts) | Fast iteration, UI-first         | Low determinism, shallow automation | Low                  |
| Figma (Weave)      | Design + proto-AI     | Layout, thumbnails, UI polish        | Strong design system alignment   | Weak pipeline orchestration         | Medium (design only) |
| Fal.ai (Workflows) | Backend orchestration | Model chaining (video, audio)        | Scalable, API-native             | Requires infra thinking             | High                 |
| Flora              | Generative canvas     | Concept visuals                      | Fast creative exploration        | Not production-oriented             | Low                  |
| InvokeAI           | Node-based (images)   | Controlled image gen                 | Deterministic, local             | Weak for video                      | Medium (thumbnails)  |
| Playbook3D         | 3D asset gen          | Scenes / environments                | Structured spatial control       | Overkill, slow                      | Low                  |
| ComfyUI            | Full graph control    | Image/video pipelines                | Maximum control, reproducibility | Complexity, setup cost              | Very High            |
| n8n                | Automation            | Glue (scripts → TTS → gen → export)  | Deterministic workflows          | No media-native ops                 | Very High            |

Functional Mapping to Your Pipeline

| Stage                 | Best tool        | Why                    |
|-----------------------|------------------|------------------------|
| Hook/script gen       | n8n + LLM        | automation + branching |
| Voice gen             | n8n → ElevenLabs | API chaining           |
| Scene generation      | ComfyUI / Fal    | control vs scale       |
| Asset iteration       | ComfyUI          | reproducibility        |
| Thumbnail design      | Figma / Invoke   | precision              |
| Orchestration         | n8n / Fal        | deterministic          |
| Experimental ideation | Krea / Flora     | fast exploration       |

Key Tradeoffs

| Axis            | ComfyUI     | Fal Workflows | n8n      |
|-----------------|-------------|---------------|----------|
| Control         | Max         | High          | Low      |
| Scalability     | Local-bound | Cloud-scale   | External |
| Determinism     | High        | Medium        | High     |
| Setup cost      | High        | Medium        | Low      |
| Media awareness | High        | Medium        | None     |

+ Ask anything

↓